

```

import pandas as pd
import numpy as np
import seaborn as sns
import os
from PIL import Image, ImageOps
from sklearn.model_selection import train_test_split

from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Activation, Dropout, Flatten, Dense
from keras import optimizers
from keras.preprocessing.image import ImageDataGenerator
import tensorflow as tf

```

## 2. Loading Data

```

In [2]:
images = []
ages = []
genders = []

for i in os.listdir('../input/utkface-new/crop_part1/')[0:8000]:
    split = i.split('_')
    ages.append(int(split[0]))
    genders.append(int(split[1]))
    images.append(Image.open('../input/utkface-new/crop_part1/' + i))

```

```

In [3]:
images = pd.Series(list(images), name = 'Images')
ages = pd.Series(list(ages), name = 'Ages')
genders = pd.Series(list(genders), name = 'Genders')

df = pd.concat([images, ages, genders], axis=1)
df

```

Out[3]:

	Images	Ages	Genders
0	<PIL.JpegImagePlugin.JpegImageFile image mode=...	26	0
1	<PIL.JpegImagePlugin.JpegImageFile image mode=...	21	1
2	<PIL.JpegImagePlugin.JpegImageFile image mode=...	17	1
3	<PIL.JpegImagePlugin.JpegImageFile image mode=...	76	0
4	<PIL.JpegImagePlugin.JpegImageFile image mode=...	18	1
...	...	...	...
7995	<PIL.JpegImagePlugin.JpegImageFile image mode=...	3	0
7996	<PIL.JpegImagePlugin.JpegImageFile image mode=...	28	0
7997	<PIL.JpegImagePlugin.JpegImageFile image mode=...	10	0
7998	<PIL.JpegImagePlugin.JpegImageFile image mode=...	8	1
7999	<PIL.JpegImagePlugin.JpegImageFile image mode=...	22	0

8000 rows × 3 columns

```
In [4]: display(df['Images'][0])
print(df['Ages'][0], df['Genders'][0])
```



26 0

```
In [5]: display(df['Images'][1])
print(df['Ages'][1], df['Genders'][1])
```



21 1

So 0 corresponds to male, 1 corresponds to female.


### 3. Visualising and Preparing Data

```
In [6]: sns.set_theme()
sns.distplot(df['Ages'], kde=True, bins=30)
```

/opt/conda/lib/python3.7/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[6]: <AxesSubplot:xlabel='Ages', ylabel='Density'>
```



Too many faces of people between 0 and 4 years old. The model would fit too well to these ages and not enough to the other ages. To resolve this I'm only going to include a third of the images between these ages.

```
In [7]:
under4s = []

for i in range(len(df)):
    if df['Ages'].iloc[i] <= 4:
        under4s.append(df.iloc[i])
under4s = pd.DataFrame(under4s)
under4s = under4s.sample(frac=0.3)

df = df[df['Ages'] > 4]


df = pd.concat([df, under4s], ignore_index = True)
```

```
In [8]:
sns.distplot(df['Ages'],kde=True, bins=30)
```

/opt/conda/lib/python3.7/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
Out[8]:
<AxesSubplot:xlabel='Ages', ylabel='Density'>
```



This looks much better! The dataframe is more representative of the population now. However, there aren't many images of people over 80, which would cause the model to not train well enough on those ages. It's best to just remove over 80s and only have a model that can predict the ages of people under 80.

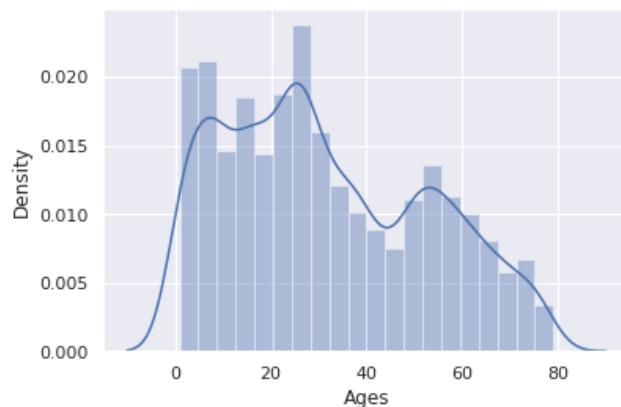
```
In [9]: df = df[df['Ages'] < 80]
```

```
In [10]: sns.distplot(df['Ages'], kde=True, bins=20)
```

```
/opt/conda/lib/python3.7/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
warnings.warn(msg, FutureWarning)
```

```
Out[10]: <AxesSubplot:xlabel='Ages', ylabel='Density'>
```




```
In [11]: sns.countplot(df['Genders'])
```

```
/opt/conda/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
```

```
Out[11]: <AxesSubplot:xlabel='Genders', ylabel='count'>
```



Not sure what 3 corresponds to - both genders, no gender, unknown, or just an error in the data entry? To be safe, I am going to remove any rows where gender equals 3.

In [12]:


```
df = df[df['Genders'] != 3]
sns.countplot(df['Genders'])
```

/opt/conda/lib/python3.7/site-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[12]:

<AxesSubplot:xlabel='Genders', ylabel='count'>



In [13]:

```
x = []
y = []

for i in range(len(df)):
    df['Images'].iloc[i] = df['Images'].iloc[i].resize((200,200), Image.ANTIALIAS)
    ar = np.asarray(df['Images'].iloc[i])
    x.append(ar)
    agegen = [int(df['Ages'].iloc[i]), int(df['Genders'].iloc[i])]
    y.append(agegen)
x = np.array(x)
```

```
/opt/conda/lib/python3.7/site-packages/pandas/core/indexing.py:670: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
`iloc._setitem_with_indexer(indexer, value)`

## 4. Train Test Split

In [14]:

```
y_age = df['Ages']  
y_gender = df['Genders']  
  
x_train_age, x_test_age, y_train_age, y_test_age = train_test_split(x, y_age, test_size=0.2, stratify=y_age)  
x_train_gender, x_test_gender, y_train_gender, y_test_gender = train_test_split(x, y_gender, test_size=0.2, stratify=y_gender)
```

## 5. Creating the Models

I will create two individual models - one to predict age and one to predict gender. The age model should be capable of returning continuous values which I will round to the nearest integer, and the gender model should return a binary result.

In [15]:

```
agemodel = Sequential()
agemodel.add(Conv2D(32, (3,3), activation='relu', input_shape=(200, 200, 3)))
agemodel.add(MaxPooling2D((2,2)))
agemodel.add(Conv2D(64, (3,3), activation='relu'))
agemodel.add(MaxPooling2D((2,2)))
agemodel.add(Conv2D(128, (3,3), activation='relu'))
agemodel.add(MaxPooling2D((2,2)))
agemodel.add(Flatten())
agemodel.add(Dense(64, activation='relu'))
agemodel.add(Dropout(0.5))
agemodel.add(Dense(1, activation='relu'))

agemodel.compile(loss='mean_squared_error',
                  optimizer=optimizers.Adam(lr=0.0001))

genmodel = Sequential()
genmodel.add(Conv2D(32, (3,3), activation='relu', input_shape=(200, 200, 3)))
genmodel.add(MaxPooling2D((2,2)))
genmodel.add(Conv2D(64, (3,3), activation='relu'))
genmodel.add(MaxPooling2D((2,2)))
genmodel.add(Conv2D(128, (3,3), activation='relu'))
genmodel.add(MaxPooling2D((2,2)))
genmodel.add(Flatten())
genmodel.add(Dense(64, activation='relu'))
genmodel.add(Dropout(0.5))
genmodel.add(Dense(1, activation='sigmoid'))

genmodel.compile(loss='binary_crossentropy',
                  optimizer=optimizers.Adam(lr=0.0001),
                  metrics=['accuracy'])
```

## 6. Training the Models

In [16]:

```
datagen = ImageDataGenerator(
    rescale=1./255., width_shift_range = 0.1, height_shift_range = 0.1, horizontal_flip = True)

test_datagen = ImageDataGenerator(rescale=1./255)

train1 = datagen.flow(x_train_age, y_train_age, batch_size=32)

test1 = test_datagen.flow(
    x_test_age, y_test_age,
    batch_size=32)

history1 = agemodel.fit(train1, epochs=50, shuffle=True, validation_data=test1)
```

Epoch 1/50  
161/161 [=====] - 44s 253ms/step - loss: 706.2857 - val\_loss: 452.0310  
Epoch 2/50  
161/161 [=====] - 41s 255ms/step - loss: 493.1964 - val\_loss: 403.0540  
Epoch 3/50  
161/161 [=====] - 42s 262ms/step - loss: 439.5375 - val\_loss: 419.8955  
Epoch 4/50  
161/161 [=====] - 41s 257ms/step - loss: 382.1004 - val\_loss: 328.3317  
Epoch 5/50  
161/161 [=====] - 41s 256ms/step - loss: 376.2793 - val\_loss: 306.8690  
Epoch 6/50  
161/161 [=====] - 41s 255ms/step - loss: 378.7557 - val\_loss: 298.7830  
Epoch 7/50  
161/161 [=====] - 42s 258ms/step - loss: 353.1836 - val\_loss: 364.5480  
Epoch 8/50  
161/161 [=====] - 41s 257ms/step - loss: 350.4301 - val\_loss: 273.9537  
Epoch 9/50  
161/161 [=====] - 41s 257ms/step - loss: 347.8215 - val\_loss: 281.0771  
Epoch 10/50  
161/161 [=====] - 41s 257ms/step - loss: 316.9603 - val\_loss: 322.0839  
Epoch 11/50  
161/161 [=====] - 42s 260ms/step - loss: 321.0114 - val\_loss: 257.4635  
Epoch 12/50  
161/161 [=====] - 42s 259ms/step - loss: 307.2590 - val\_loss: 299.9180  
Epoch 13/50  
161/161 [=====] - 41s 256ms/step - loss: 308.6267 - val\_loss: 329.4791  
Epoch 14/50  
161/161 [=====] - 42s 261ms/step - loss: 314.8022 - val\_loss: 245.9729  
Epoch 15/50  
161/161 [=====] - 41s 257ms/step - loss: 301.3866 - val\_loss: 255.6889  
Epoch 16/50  
161/161 [=====] - 41s 258ms/step - loss: 287.3217 - val\_loss: 238.2000  
Epoch 17/50  
161/161 [=====] - 41s 258ms/step - loss: 299.2067 - val\_loss: 226.4820  
Epoch 18/50  
161/161 [=====] - 42s 261ms/step - loss: 293.4950 - val\_loss: 219.9284  
Epoch 19/50  
161/161 [=====] - 41s 258ms/step - loss: 268.2615 - val\_loss: 218.2262  
Epoch 20/50  
161/161 [=====] - 41s 254ms/step - loss: 284.8218 - val\_loss: 234.8680  
Epoch 21/50  
161/161 [=====] - 41s 258ms/step - loss: 267.3226 - val\_loss: 216.5874  
Epoch 22/50  
161/161 [=====] - 42s 260ms/step - loss: 280.1753 - val\_loss: 238.4909  
Epoch 23/50  
161/161 [=====] - 41s 257ms/step - loss: 269.4055 - val\_loss: 200.0749  
Epoch 24/50  
161/161 [=====] - 41s 256ms/step - loss: 262.5218 - val\_loss: 221.6551  
Epoch 25/50  
161/161 [=====] - 42s 260ms/step - loss: 275.6070 - val\_loss: 198.7766  
Epoch 26/50  
161/161 [=====] - 42s 259ms/step - loss: 266.5272 - val\_loss: 221.8195  
Epoch 27/50  
161/161 [=====] - 41s 257ms/step - loss: 254.4120 - val\_loss: 199.4566  
Epoch 28/50



161/161 [=====] - 41s 256ms/step - loss: 252.5054 - val\_loss: 193.0954  
Epoch 29/50  
161/161 [=====] - 42s 259ms/step - loss: 267.9600 - val\_loss: 196.8834  
Epoch 30/50  
161/161 [=====] - 41s 258ms/step - loss: 241.8575 - val\_loss: 187.8295  
Epoch 31/50  
161/161 [=====] - 41s 256ms/step - loss: 241.0621 - val\_loss: 183.0156  
Epoch 32/50  
161/161 [=====] - 42s 259ms/step - loss: 238.2132 - val\_loss: 189.9071  
Epoch 33/50  
161/161 [=====] - 41s 256ms/step - loss: 238.4509 - val\_loss: 193.3439  
Epoch 34/50  
161/161 [=====] - 41s 257ms/step - loss: 230.6079 - val\_loss: 197.5394  
Epoch 35/50  
161/161 [=====] - 41s 255ms/step - loss: 252.5103 - val\_loss: 177.1196  
Epoch 36/50  
161/161 [=====] - 42s 260ms/step - loss: 248.8367 - val\_loss: 190.6460  
Epoch 37/50  
161/161 [=====] - 41s 254ms/step - loss: 238.1350 - val\_loss: 222.1548  
Epoch 38/50  
161/161 [=====] - 41s 257ms/step - loss: 232.1687 - val\_loss: 171.6752  
Epoch 39/50  
161/161 [=====] - 41s 257ms/step - loss: 227.3713 - val\_loss: 180.7836  
Epoch 40/50  
161/161 [=====] - 42s 260ms/step - loss: 229.9505 - val\_loss: 167.7247  
Epoch 41/50  
161/161 [=====] - 41s 257ms/step - loss: 229.1870 - val\_loss: 221.3032  
Epoch 42/50  
161/161 [=====] - 42s 257ms/step - loss: 232.8839 - val\_loss: 166.4539  
Epoch 43/50  
161/161 [=====] - 41s 255ms/step - loss: 228.9473 - val\_loss: 161.2658  
Epoch 44/50  
161/161 [=====] - 42s 260ms/step - loss: 222.0279 - val\_loss: 185.0128  
Epoch 45/50  
161/161 [=====] - 41s 256ms/step - loss: 219.0961 - val\_loss: 160.4378  
Epoch 46/50  
161/161 [=====] - 41s 253ms/step - loss: 221.8582 - val\_loss: 174.0324  
Epoch 47/50  
161/161 [=====] - 41s 256ms/step - loss: 212.9477 - val\_loss: 159.7157  
Epoch 48/50  
161/161 [=====] - 42s 259ms/step - loss: 220.9071 - val\_loss: 189.6833  
Epoch 49/50  
161/161 [=====] - 42s 257ms/step - loss: 232.4877 - val\_loss: 160.9419  
Epoch 50/50  
161/161 [=====] - 41s 253ms/step - loss: 212.5144 - val\_loss: 153.1291

In [17]:

```
datagen = ImageDataGenerator(  
    rescale=1./255., width_shift_range = 0.1, height_shift_range = 0.1, horizontal_flip = True)  
  
test_datagen = ImageDataGenerator(rescale=1./255)  
  
train2 = datagen.flow(x_train_gender, y_train_gender, batch_size=64)  
  
test2 = test_datagen.flow(  
    x_test_gender, y_test_gender,  
    batch_size=64)  
  
history2 = genmodel.fit(train2, epochs=50, shuffle=True, validation_data=test2)
```

Epoch 1/50  
81/81 [=====] - 41s 489ms/step - loss: 0.6716 - accuracy: 0.5686 - val\_loss: 0.6030 - val\_accuracy: 0.6568  
Epoch 2/50  
81/81 [=====] - 40s 492ms/step - loss: 0.6089 - accuracy: 0.6726 - val\_loss: 0.5504 - val\_accuracy: 0.7189  
Epoch 3/50  
81/81 [=====] - 40s 492ms/step - loss: 0.5818 - accuracy: 0.6943 - val\_loss: 0.5240 - val\_accuracy: 0.7399  
Epoch 4/50  
81/81 [=====] - 40s 496ms/step - loss: 0.5557 - accuracy: 0.7236 - val\_loss: 0.4831 - val\_accuracy: 0.7896  
Epoch 5/50  
81/81 [=====] - 40s 498ms/step - loss: 0.5302 - accuracy: 0.7449 - val\_loss: 0.4828 - val\_accuracy: 0.7857  
Epoch 6/50  
81/81 [=====] - 40s 496ms/step - loss: 0.4865 - accuracy: 0.7740 - val\_loss: 0.4606 - val\_accuracy: 0.8028  
Epoch 7/50  
81/81 [=====] - 41s 506ms/step - loss: 0.4879 - accuracy: 0.7628 - val\_loss: 0.4375 - val\_accuracy: 0.8005  
Epoch 8/50  
81/81 [=====] - 40s 495ms/step - loss: 0.4751 - accuracy: 0.7678 - val\_loss: 0.4417 - val\_accuracy: 0.8051  
Epoch 9/50  
81/81 [=====] - 40s 492ms/step - loss: 0.4626 - accuracy: 0.7854 - val\_loss: 0.4303 - val\_accuracy: 0.8005  
Epoch 10/50  
81/81 [=====] - 40s 497ms/step - loss: 0.4556 - accuracy: 0.7767 - val\_loss: 0.4127 - val\_accuracy: 0.8113  
Epoch 11/50  
81/81 [=====] - 41s 507ms/step - loss: 0.4504 - accuracy: 0.7882 - val\_loss: 0.4207 - val\_accuracy: 0.8051  
Epoch 12/50  
81/81 [=====] - 40s 490ms/step - loss: 0.4513 - accuracy: 0.7927 - val\_loss: 0.4068 - val\_accuracy: 0.8129  
Epoch 13/50  
81/81 [=====] - 40s 500ms/step - loss: 0.4346 - accuracy: 0.8021 - val\_loss: 0.3897 - val\_accuracy: 0.8168  
Epoch 14/50  
81/81 [=====] - 40s 498ms/step - loss: 0.4434 - accuracy: 0.7828 - val\_loss: 0.3890 - val\_accuracy: 0.8269  
Epoch 15/50  
81/81 [=====] - 40s 494ms/step - loss: 0.4175 - accuracy: 0.8103 - val\_loss: 0.3905 - val\_accuracy: 0.8207  
Epoch 16/50  
81/81 [=====] - 41s 502ms/step - loss: 0.4232 - accuracy: 0.8068 - val\_loss: 0.3901 - val\_accuracy: 0.8261  
Epoch 17/50  
81/81 [=====] - 40s 498ms/step - loss: 0.4193 - accuracy: 0.8055 - val\_loss: 0.3899 - val\_accuracy: 0.8261  
Epoch 18/50  
81/81 [=====] - 40s 496ms/step - loss: 0.4226 - accuracy: 0.7830 - val\_loss: 0.3832 - val\_accuracy: 0.8276  
Epoch 19/50

81/81 [=====] - 40s 493ms/step - loss: 0.4147 - accuracy: 0.8058 - val\_loss: 0.3842 - val\_accuracy: 0.8276  
Epoch 20/50  
81/81 [=====] - 41s 501ms/step - loss: 0.4129 - accuracy: 0.8096 - val\_loss: 0.3787 - val\_accuracy: 0.8276  
Epoch 21/50  
81/81 [=====] - 40s 496ms/step - loss: 0.3959 - accuracy: 0.8106 - val\_loss: 0.3689 - val\_accuracy: 0.8346  
Epoch 22/50  
81/81 [=====] - 40s 500ms/step - loss: 0.4189 - accuracy: 0.8021 - val\_loss: 0.3724 - val\_accuracy: 0.8315  
Epoch 23/50  
81/81 [=====] - 40s 495ms/step - loss: 0.4055 - accuracy: 0.8033 - val\_loss: 0.3785 - val\_accuracy: 0.8284  
Epoch 24/50  
81/81 [=====] - 40s 500ms/step - loss: 0.3869 - accuracy: 0.8192 - val\_loss: 0.3640 - val\_accuracy: 0.8292  
Epoch 25/50  
81/81 [=====] - 40s 495ms/step - loss: 0.3761 - accuracy: 0.8286 - val\_loss: 0.3706 - val\_accuracy: 0.8346  
Epoch 26/50  
81/81 [=====] - 40s 492ms/step - loss: 0.3827 - accuracy: 0.8313 - val\_loss: 0.3658 - val\_accuracy: 0.8362  
Epoch 27/50  
81/81 [=====] - 40s 492ms/step - loss: 0.4041 - accuracy: 0.8122 - val\_loss: 0.3666 - val\_accuracy: 0.8424  
Epoch 28/50  
81/81 [=====] - 41s 502ms/step - loss: 0.3957 - accuracy: 0.8220 - val\_loss: 0.3565 - val\_accuracy: 0.8416  
Epoch 29/50  
81/81 [=====] - 41s 501ms/step - loss: 0.3940 - accuracy: 0.8238 - val\_loss: 0.3803 - val\_accuracy: 0.8269  
Epoch 30/50  
81/81 [=====] - 40s 490ms/step - loss: 0.3837 - accuracy: 0.8195 - val\_loss: 0.3705 - val\_accuracy: 0.8377  
Epoch 31/50  
81/81 [=====] - 40s 498ms/step - loss: 0.3899 - accuracy: 0.8177 - val\_loss: 0.3597 - val\_accuracy: 0.8385  
Epoch 32/50  
81/81 [=====] - 41s 506ms/step - loss: 0.3966 - accuracy: 0.8177 - val\_loss: 0.3592 - val\_accuracy: 0.8393  
Epoch 33/50  
81/81 [=====] - 40s 493ms/step - loss: 0.3782 - accuracy: 0.8214 - val\_loss: 0.3543 - val\_accuracy: 0.8494  
Epoch 34/50  
81/81 [=====] - 40s 495ms/step - loss: 0.3896 - accuracy: 0.8154 - val\_loss: 0.3470 - val\_accuracy: 0.8463  
Epoch 35/50  
81/81 [=====] - 41s 501ms/step - loss: 0.3743 - accuracy: 0.8232 - val\_loss: 0.3432 - val\_accuracy: 0.8478  
Epoch 36/50  
81/81 [=====] - 41s 500ms/step - loss: 0.3786 - accuracy: 0.8239 - val\_loss: 0.3402 - val\_accuracy: 0.8455  
Epoch 37/50  
81/81 [=====] - 40s 499ms/step - loss: 0.3554 - accuracy: 0.8404 - val\_loss: 0.3512 - val\_accuracy: 0.8463

Epoch 38/50  
81/81 [=====] - 40s 496ms/step - loss: 0.3809 - accuracy: 0.8286 - val\_loss: 0.3491 - val\_accuracy: 0.8540  
Epoch 39/50  
81/81 [=====] - 40s 492ms/step - loss: 0.3658 - accuracy: 0.8344 - val\_loss: 0.3468 - val\_accuracy: 0.8463  
Epoch 40/50  
81/81 [=====] - 41s 505ms/step - loss: 0.3718 - accuracy: 0.8291 - val\_loss: 0.3395 - val\_accuracy: 0.8509  
Epoch 41/50  
81/81 [=====] - 40s 496ms/step - loss: 0.3642 - accuracy: 0.8309 - val\_loss: 0.3519 - val\_accuracy: 0.8494  
Epoch 42/50  
81/81 [=====] - 40s 489ms/step - loss: 0.3658 - accuracy: 0.8329 - val\_loss: 0.3399 - val\_accuracy: 0.8447  
Epoch 43/50  
81/81 [=====] - 40s 496ms/step - loss: 0.3632 - accuracy: 0.8328 - val\_loss: 0.3407 - val\_accuracy: 0.8447  
Epoch 44/50  
81/81 [=====] - 41s 506ms/step - loss: 0.3569 - accuracy: 0.8474 - val\_loss: 0.3459 - val\_accuracy: 0.8502  
Epoch 45/50  
81/81 [=====] - 40s 495ms/step - loss: 0.3489 - accuracy: 0.8388 - val\_loss: 0.3409 - val\_accuracy: 0.8463  
Epoch 46/50  
81/81 [=====] - 40s 495ms/step - loss: 0.3607 - accuracy: 0.8433 - val\_loss: 0.3385 - val\_accuracy: 0.8494  
Epoch 47/50  
81/81 [=====] - 40s 489ms/step - loss: 0.3515 - accuracy: 0.8444 - val\_loss: 0.3429 - val\_accuracy: 0.8502  
Epoch 48/50  
81/81 [=====] - 40s 497ms/step - loss: 0.3520 - accuracy: 0.8401 - val\_loss: 0.3361 - val\_accuracy: 0.8556  
Epoch 49/50  
81/81 [=====] - 40s 496ms/step - loss: 0.3475 - accuracy: 0.8469 - val\_loss: 0.3392 - val\_accuracy: 0.8470  
Epoch 50/50  
81/81 [=====] - 40s 495ms/step - loss: 0.3502 - accuracy: 0.8380 - val\_loss: 0.3420 - val\_accuracy: 0.8564

## 7. Evaluating the Models

Now to evaluate the models I am going to use some external images of celebrities. These celebrities are of a variety of ages and genders.

In [30]:

```
def process_and_predict(file):
    im = Image.open(file)
    width, height = im.size
    if width == height:
        im = im.resize((200,200), Image.ANTIALIAS)
    else:
        if width > height:
            left = width/2 - height/2
            right = width/2 + height/2
            top = 0
            bottom = height
            im = im.crop((left,top,right,bottom))
            im = im.resize((200,200), Image.ANTIALIAS)
        else:
            left = 0
            right = width
            top = 0
            bottom = width
            im = im.crop((left,top,right,bottom))
            im = im.resize((200,200), Image.ANTIALIAS)

    ar = np.asarray(im)
    ar = ar.astype('float32')
    ar /= 255.0
    ar = ar.reshape(-1, 200, 200, 3)

    age = agemodel.predict(ar)
    gender = np.round(genmodel.predict(ar))
    if gender == 0:
        gender = 'male'
    elif gender == 1:
        gender = 'female'

    print('Age:', int(age), '\n Gender:', gender)
    return im.resize((300,300), Image.ANTIALIAS)
```

## Alyson Hannigan

In [43]:

```
process_and_predict('../input/celebrities2/alyson.jpg')
```

Age: 36

Gender: female

Out[43]:



David Boreanaz

In [46]:

```
process_and_predict('../input/celebrities2/david.jpg')
```

Age: 61

Gender: male

Out[46]:



Gaten Matarazzo

In [48]:

```
process_and_predict('../input/celebrities2/gaten.jpg')
```

Age: 38

Gender: male

Out[48]:



Jack Dylan Grazer

In [50]:

```
process_and_predict('../input/celebrities2/jack.jpg')
```

Age: 14

Gender: female

Out[50]:



Jennifer Garner

In [52]:

```
process_and_predict('../input/celebrities2/jennifer.jpg')
```

Age: 28

Gender: female



Out[52]:



Jennifer Lawrence

In [53]:

```
process_and_predict('../input/celebrities2/jenniferlaw.jpg')
```

Age: 47

Gender: female

Out[53]:



Meryl Streep

In [57]:

```
process_and_predict('../input/celebrities2/meryl.jpg')
```

Age: 64

Gender: female

Out[57]:



Millie Bobby Brown


In [58]:

```
process_and_predict('../input/celebrities2/millie.jpg')
```

Age: 13

Gender: female

Out[58]:



Morgan Freeman


In [59]:

```
process_and_predict('../input/celebrities2/morgan.jpg')
```

Age: 86

Gender: male

Out[59]:



Oprah Winfrey

In [60]:

```
process_and_predict('../input/celebrities2/oprah.jpg')
```

Age: 42

Gender: female

Out[60]:



Tom Hanks


In [63]:

```
process_and_predict('../input/celebrities2/tom.jpg')
```

Age: 51

Gender: male

Out[63]:



Winona Ryder

In [65]:

```
process_and_predict('../input/celebrities2/winona.jpg')
```

Age: 26

Gender: female

Out[65]:

